

Practical Work 3 Classification with Bayes - System Evaluation

Students

- Flavia Pittet
- Fabio Vitali

Exercise 1 Classification system using Bayes

a. Bayes - Histograms

Implement a classifier based on Bayes using histograms to estimate the likelihoods.

- a) Read the training data from file `ex1-data-train.csv`. The first two columns are `x1` and `x2`. The last column holds the class label `y`.
- b) Compute the priors of both classes $P(C0)$ and $P(C1)$.
- c) Compute histograms of `x1` and `x2` for each class (total of 4 histograms). Plot these histograms. Advice : use the `numpy histogram(a,bins='auto')` function.

In [2]:

```
# basic imports
import pandas as pd
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import numpy as np
import random
from mpl_toolkits.mplot3d import Axes3D
%matplotlib inline

# a) read data
trainingset = pd.read_csv('data/ex1-data-train.csv', names=['grade1', 'grade2',
'pass'])
x1 = trainingset['grade1'].values
x2 = trainingset['grade2'].values
y = trainingset['pass'].values
trainset = list(zip(x1, x2, y))

testset = pd.read_csv('data/ex1-data-test.csv', names=['grade1', 'grade2', 'pas
s'])
testx1 = testset['grade1'].values
testx2 = testset['grade2'].values
testy = testset['pass'].values

testset = list(zip(testx1, testx2, testy))

## b) compute priors
priorCpass = sum([1 for x1, x2, y in trainset if y])/len(trainset)

priorCfail = 1 - priorCpass

print('pass rate = {} \n fail rate = {}'.format(priorCpass, priorCfail))
## c) histogram
#[x for x,_,_ in trainset]
#[y for _,_,y in trainset]
fig = plt.figure(figsize=(15,15))

pltx1Success = fig.add_subplot(2, 2, 1)
x1Success = [x for x,_,y in trainset if y]
x1SuccessValues, x1SuccessBins = np.histogram(x1Success, bins='auto')
pltx1Success.hist(x=x1Success, bins=x1SuccessBins, color='green')
pltx1Success.set_title('x1 Grade Success')

pltx1Fail = fig.add_subplot(2, 2, 2)
x1Fail = [x for x,_,y in trainset if not y]
x1FailValues, x1FailBins = np.histogram(x1Success, bins='auto')
pltx1Fail.hist(x=x1Fail, bins=x1FailBins, color='red')
pltx1Fail.set_title('x1 Grade Fail')

print(x1SuccessValues, x1SuccessBins)

pltx2Success = fig.add_subplot(2, 2, 3)
x2Success = [x for _,x,y in trainset if y]
x2SuccessValues, x2SuccessBins = np.histogram(x2Success, bins='auto')
pltx2Success.hist(x=x2Success, bins=x2SuccessBins, color='green')
pltx2Success.set_title('x2 Grade Success')

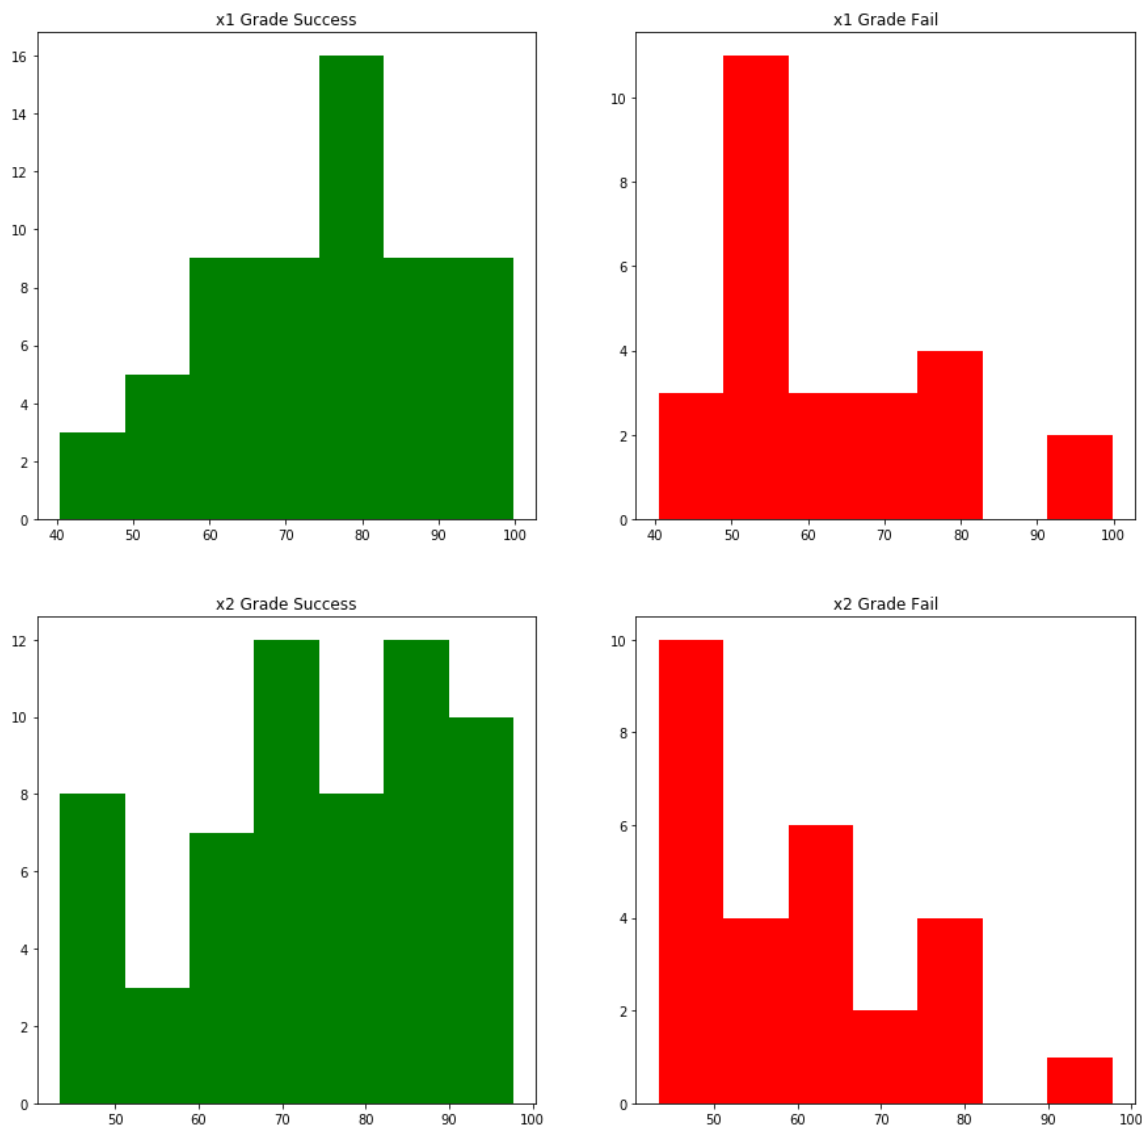
pltx2Fail = fig.add_subplot(2, 2, 4)
```

```
x2Fail = [x for _,x,y in trainset if not y]
x2FailValues, x2FailBins = np.histogram(x2Success, bins='auto')
plt.x2Fail.hist(x=x2Fail, bins=x2FailBins, color='red')
plt.x2Fail.set_title('x2 Grade Fail')
```

```
pass rate = 0.6
fail rate = 0.4
[ 3  5  9  9 16  9  9] [40.45755098 48.93902339 57.42049579 65.90196
819 74.38344059 82.86491299
91.3463854 99.8278578 ]
```

Out[2]:

Text(0.5,1,'x2 Grade Fail')



d) Use the histograms to compute the likelihoods $p(x_1|C_0)$, $p(x_1|C_1)$, $p(x_2|C_0)$ and $p(x_2|C_1)$. For this define a function `likelihoodHist(x,histValues,edgeValues)` that returns the likelihood of x for a given histogram (defined by its values and bin edges as returned by the `numpy.histogram()` function).

In [3]:

```
def likelihoodHist(x,histValues,edgeValues):
    i = 0
    for i in range(len(edgeValues)):
        if edgeValues[i] >= x:
            break
    if i == 0:
        return 0

    i -= 1

    return histValues[i] / sum(histValues)
```

e) Implement the classification decision according to Bayes rule and compute the overall accuracy of the system on the test set `ex1-data-test.csv`. :

- using only feature `x1`
- using only feature `x2`
- using `x1` and `x2` making the naive Bayes hypothesis of feature independence, i.e. $p(X|C_k) = p(x_1|C_k) \cdot p(x_2|C_k)$

Which system is the best ?

The last one because it takes more parameters into consideration and therefore has a better success rate.

In [4]:

```
def bayes_rule_x1(x):
    likelihoodSuccess = likelihoodHist(x, x1SuccessValues, x1SuccessBins) * priorCpass
    likelihoodFail = likelihoodHist(x, x1FailValues, x1FailBins) * priorCfail
    return likelihoodSuccess > likelihoodFail

def bayes_rule_x2(x):
    likelihoodSuccess = likelihoodHist(x, x2SuccessValues, x2SuccessBins) * priorCpass
    likelihoodFail = likelihoodHist(x, x2FailValues, x2FailBins) * priorCfail
    return likelihoodSuccess > likelihoodFail

def bayes_rule(x1, x2):
    likelihoodSuccessX1 = likelihoodHist(x1, x1SuccessValues, x1SuccessBins) * priorCpass
    likelihoodFailX1 = likelihoodHist(x1, x1FailValues, x1FailBins) * priorCfail
    likelihoodSuccessX2 = likelihoodHist(x2, x2SuccessValues, x2SuccessBins) * priorCpass
    likelihoodFailX2 = likelihoodHist(x2, x2FailValues, x2FailBins) * priorCfail

    return likelihoodSuccessX1 * likelihoodSuccessX2 > likelihoodFailX1 * likelihoodFailX2

successX1 = 0
successX2 = 0
successBoth = 0
for x1, x2, y in testset:

    if bayes_rule_x1(x1) == y:
        successX1 += 1
    if bayes_rule_x2(x2) == y:
        successX2 += 1
    if bayes_rule(x1, x2) == y:
        successBoth += 1

print('Success rate bayes on x1 = {}'.format(successX1/len(testset)))
print('Success rate bayes on x2 = {}'.format(successX2/len(testset)))
print('Success rate bayes on x1 and x2 = {}'.format(successBoth/len(testset)))
```

```
Success rate bayes on x1 = 0.7
Success rate bayes on x2 = 0.67
Success rate bayes on x1 and x2 = 0.85
```

b. Bayes - Univariate Gaussian distribution

Do the same as in c. but this time using univariate Gaussian distribution to model the likelihoods $p(x_1|C_0)$, $p(x_1|C_1)$, $p(x_2|C_0)$ and $p(x_2|C_1)$. You may use the numpy functions `mean()` and `var()` to compute the mean μ and variance σ^2 of the distribution. To model the likelihood of both features, you may also do the naive Bayes hypothesis of feature independence, i.e. $p(X|C_k) = p(x_1|C_k) \cdot p(x_2|C_k)$.

In [5]:

```
x1SuccessMean = np.mean(x1Success)
x1SuccessVar = np.var(x1Success)

x1FailMean = np.mean(x1Fail)
x1FailVar = np.var(x1Fail)

x2SuccessMean = np.mean(x2Success)
x2SuccessVar = np.var(x2Success)

x2FailMean = np.mean(x2Fail)
x2FailVar = np.var(x2Fail)

def univariate_gaussian_distribution(x, mean, var):
    v1 = (2*np.pi*var)**0.5
    v2 = np.exp(-(x-mean)**2 / (2*var))

    return v2/v1

fig = plt.figure(figsize=(15,15))

pltX1Success = fig.add_subplot(2, 2, 1)
pltX1Success.set_title('X1 Success')
pltX1Success.plot(range(100), [univariate_gaussian_distribution(x, x1SuccessMean
, x1SuccessVar) for x in range(100)], color='green')

pltX1Fail = fig.add_subplot(2, 2, 2)
pltX1Fail.set_title('X1 Fail')
pltX1Fail.plot(range(100), [univariate_gaussian_distribution(x, x1FailMean, x1Fa
ilVar) for x in range(100)], color='red')

pltX2Success = fig.add_subplot(2, 2, 3)
pltX2Success.set_title('X2 Success')
pltX2Success.plot(range(100), [univariate_gaussian_distribution(x, x2SuccessMean
, x2SuccessVar) for x in range(100)], color='green')

pltX2Fail = fig.add_subplot(2, 2, 4)
pltX2Fail.set_title('X2 Fail')
pltX2Fail.plot(range(100), [univariate_gaussian_distribution(x, x2FailMean, x2Fa
ilVar) for x in range(100)], color='red')

def univariate_gaussian_distribution_x1(x):
    likelihoodSuccess = univariate_gaussian_distribution(x, x1SuccessMean, x1Suc
cessVar) * priorCpass
    likelihoodFail = univariate_gaussian_distribution(x, x1FailMean, x1Fai
lVar) * priorCfail

    return likelihoodSuccess > likelihoodFail

def univariate_gaussian_distribution_x2(x):
    likelihoodSuccess = univariate_gaussian_distribution(x, x2SuccessMean, x2Suc
cessVar) * priorCpass
    likelihoodFail = univariate_gaussian_distribution(x, x2FailMean, x2Fai
lVar) * priorCfail

    return likelihoodSuccess > likelihoodFail
```

```

def univariate_gaussian_distribution_both(x1, x2):
    likelihoodSuccessx1 = univariate_gaussian_distribution(x1, x1SuccessMean, x1
SuccessVar) * priorCpass
    likelihoodFailx1     = univariate_gaussian_distribution(x1, x1FailMean,    x1
FailVar)      * priorCfail

    likelihoodSuccessx2 = univariate_gaussian_distribution(x2, x2SuccessMean, x2
SuccessVar) * priorCpass
    likelihoodFailx2     = univariate_gaussian_distribution(x2, x2FailMean,    x2
FailVar)      * priorCfail

    return likelihoodSuccessx1 * likelihoodSuccessx2 > likelihoodFailx1 * likeli
hoodFailx2

successX1 = 0
successX2 = 0
successBoth = 0

for x1, x2, y in testset:
    if univariate_gaussian_distribution_x1(x1) == y:
        successX1 += 1

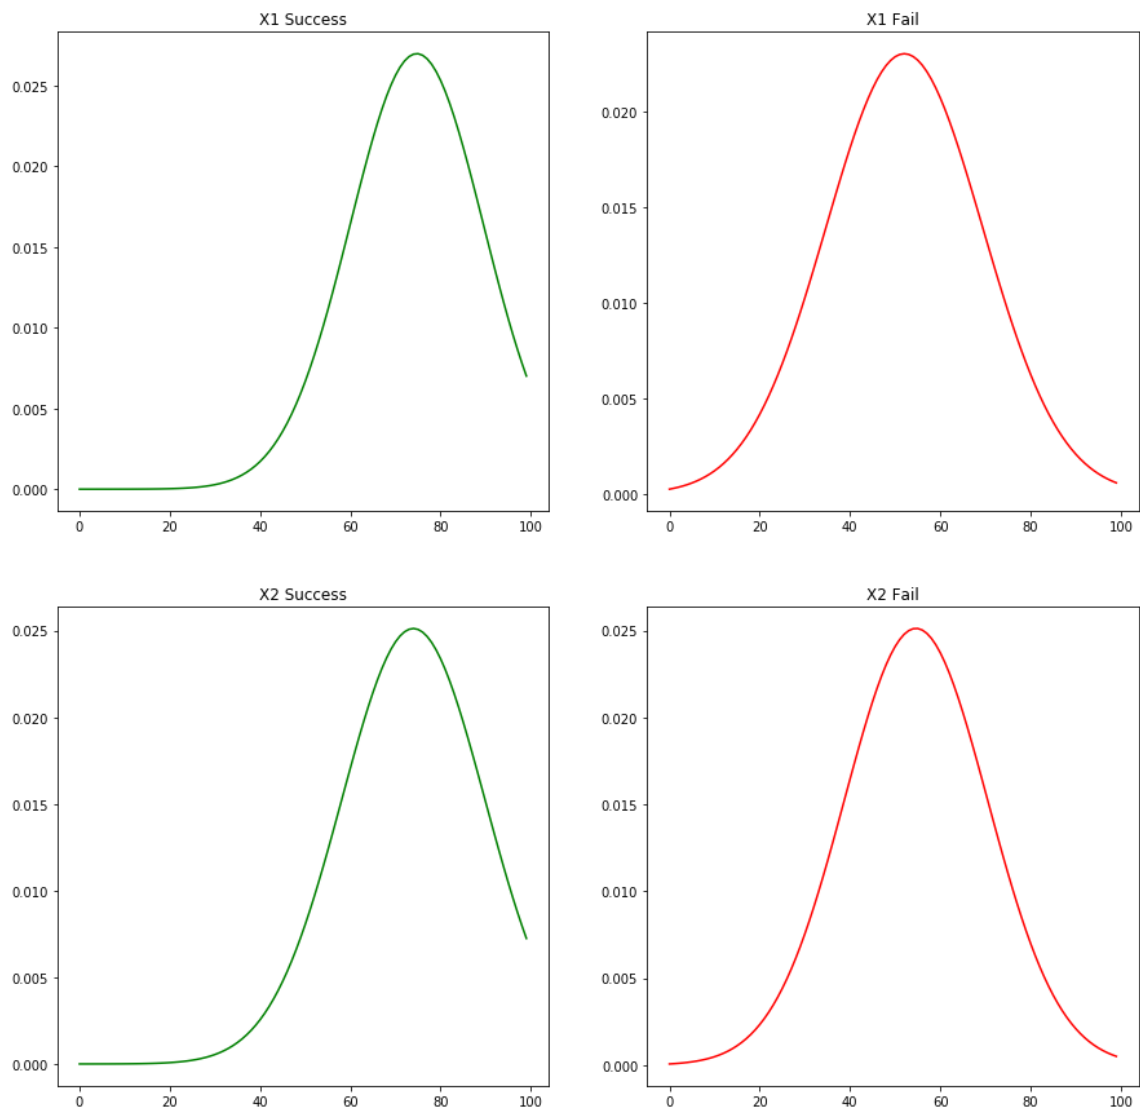
    if univariate_gaussian_distribution_x2(x2) == y:
        successX2 += 1

    if univariate_gaussian_distribution_both(x1, x2) == y:
        successBoth += 1

print('Success rate Univariate Gaussian distribution on x1          = {}'.format(s
uccessX1/len(testset)))
print('Success rate Univariate Gaussian distribution on x2          = {}'.format(s
uccessX2/len(testset)))
print('Success rate Univariate Gaussian distribution on x1 and x2 = {}'.format(s
uccessBoth/len(testset)))

```


Success rate Univariate Gaussian distribution on x1 = 0.71
Success rate Univariate Gaussian distribution on x2 = 0.72
Success rate Univariate Gaussian distribution on x1 and x2 = 0.92



Exercise 2 - System evaluation

Let's assume we have trained a digit classification system able to categorise images of digits from 0 to 9, as illustrated on Figure 2.

After training, the system has been run against a test set (independent of the training set) including $N_t = 100000$ samples. The system is able to compute estimations of a posteriori.

In this exercise, the columns are not labeled as their index is already representing the corresponding class (except for the last one which is the ground truth)

NB: as we will be using the fashion-MNIST instead of the digits, we will assume we have trained a system that can categorise these images (also from 0 to 9)

a) Write a function to take classification decisions on such outputs according to Bayes' rule

According to Bayes' rule, we need to choose the class which have the highest $P(C_k|x)$ as the predicted class

In [6]:

```
def predictionFromSystem(system):
    N = len(dataset)
    y_pred = []
    for i in range(N):
        row = predC.loc[i]
        y_hat = row[row == max(row)].index[0]
        y_pred.append(y_hat)

    return y_pred

dataset = pd.read_csv('data/ex2-system-a.csv', sep=';', header=None, usecols=[x
for x in range(11)])
predC, y_truth = dataset.iloc[:,0:10], dataset[10]

y_pred = predictionFromSystem(predC)

conf_matrix = confusion_matrix(y_truth, y_pred)
```

b) What is the overall error rate of the system ?

We compute the error directly from the confusion matrix

In [7]:

```
def computeOverallAccuracy(matrix, N):
    sum = 0
    for i in range(len(matrix)):
        sum += conf_matrix[i][i]

    return sum / N

overall_accuracy = computeOverallAccuracy(conf_matrix, len(predC))
error_rate = 1 - overall_accuracy
print("The error rate is {}".format(error_rate))
```

The error rate is 0.10729999999999995

c) Compute and report the confusion matrix of the system.

As the matrix is already computed, we just need to print it

In [8]:

```
print("Confusion Matrix:\n", conf_matrix)
```

Confusion Matrix:

```
[[ 944    0   11    0    0    2   10    7    5    1]
 [    0 1112    2    3    1    4    3    1    9    0]
 [   10    6  921   12   15    3   19   15   26    5]
 [    1    1   31  862    2   72    5   14   12   10]
 [    2    3    6    2  910    1   12    6    4   36]
 [   12    3    6   29   19  768   19    9   21    6]
 [   14    3   21    2   22   28  865    0    3    0]
 [    0   14   30    9    7    2    1  929    3   33]
 [   12   16   18   26   24   46   22   19  772   19]
 [   10    4    6   22   53   18    0   48    4  844]]
```

d) What are the worst and best classes in terms of precision and recall ?

In [9]:

```
def computeAllRecall(matrix):
    N = len(matrix)
    recall_list = [0 for _ in range(N)]
    for i in range(N):
        recall_list[i] = matrix[i][i] / matrix[i].sum()

    return recall_list

recall_list = computeAllRecall(conf_matrix)
print("The minimum recall is class {} and the maximum recall is class {}".format(
    recall_list.index(min(recall_list)),
    recall_list.index(max(recall_list))))

def computeAllClassPrecision(matrix):
    return computeAllRecall(matrix.T)

precision_list = computeAllClassPrecision(conf_matrix)
print("The minimum precision is class {} and the maximum precision is class {}".format(
    precision_list.index(min(precision_list)),
    precision_list.index(max(precision_list))))

def computeF1(rec_list, prec_list):
    F1_list = [2*(prec*recall)/(prec+recall) for recall, prec in zip(rec_list, p
rec_list)]

    return F1_list

F1_list = computeF1(recall_list, precision_list)
print("F1 for each class of System A is {}".format(F1_list))
```

The minimum recall is class 8 and the maximum recall is class 1
The minimum precision is class 5 and the maximum precision is class 1
F1 for each class of System A is [0.9511335012594457, 0.9682194166303876, 0.8838771593090212, 0.8720283257460798, 0.8943488943488943, 0.8366013071895425, 0.9038662486938349, 0.8949903660886319, 0.8423349699945445, 0.8599083036169131]

e) In file ex1-system-b.csv you find the output of a second system B. What is the best system between (a) and (b) in terms of error rate and F1

In [10]:

```
dataset = pd.read_csv('data/ex2-system-b.csv', sep=';', header=None, usecols=[x
for x in range(11)])
predC, y_truth = dataset.iloc[:,0:10], dataset[10]

y_pred = predictionFromSystem(predC)

conf_matrix = confusion_matrix(y_truth, y_pred)
```

In [11]:

```
overall_accuracy_b = computeOverallAccuracy(conf_matrix, len(predC))
error_rate_b = 1 - overall_accuracy_b
print("The error rate is {}".format(error_rate_b))

recall_list_b = computeAllRecall(conf_matrix)
precision_list_b = computeAllClassPrecision(conf_matrix)
F1_list_b = computeF1(recall_list, precision_list)

print("The minimum recall is class {} and the maximum recall is class {}".format(
    recall_list_b.index(min(recall_list_b)),

    recall_list_b.index(max(recall_list_b))))
print("The minimum precision is class {} and the maximum precision is class {}".format(
    precision_list_b.index(min(precision_list_b)),

    precision_list_b.index(max(precision_list_b))))
print("F1 for each class of System B is {}".format(F1_list_b))

print("The error rate for System A is {}, and for System B it's {}".format(error_rate, error_rate_b))
print("The difference of F1 between the 2 systems is {}".format([x-y for x,y in zip(F1_list, F1_list_b)]))
```

The error rate is 0.03869999999999996

The minimum recall is class 5 and the maximum recall is class 1

The minimum precision is class 3 and the maximum precision is class 1

F1 for each class of System B is [0.9511335012594457, 0.9682194166303876, 0.8838771593090212, 0.8720283257460798, 0.8943488943488943, 0.8366013071895425, 0.9038662486938349, 0.8949903660886319, 0.8423349699945445, 0.8599083036169131]

The error rate for System A is 0.10729999999999995, and for System B it's 0.03869999999999996

The difference of F1 between the 2 systems is [0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.0]

- We can see that System B has an error rate 3 times smaller than System A
- The F1 value we computed for each class is the same for both systems. If it's not a mistake on our part, the result seems confusing as the two systems produce the same accuracy.

Exercise 3 System evaluation

Let's look back at the PW02 exercise 3 of last week. We have built a knn classification systems for images of digits on the MNIST database.

a) How would you build a Bayesian classification for the same task ?

We have these training and test sets which have 784 features for each image (28x28 pixels)

We can construct a histogram for each possible couple (feature, class) (So 784 x 10 histograms). From them, we can compute the likelihood $P(x_n|C_k)$ for each feature. That's it for the estimator.

Now for any new image we want to test, we get its features, compute their $P(x_n|C_k)$ and use the naive Bayes hypothesis of feature independence to get $P(X|C_k) = P(x_1|C_k) \cdot P(x_2|C_k) \cdot \dots \cdot P(x_{784}|C_k)$.

Because $P(C_k)$ for any class is the same for every class, we could predict an image output simply by picking the class with the highest $P(X|C_k)$.