

07_pittet_vitali

November 5, 2018

1 Practical Work 7 - SVM

1.1 Students

- Flavia Pittet
- Fabio Vitali

1.2 Exercise - 1 Digit classification system using different SVM classifiers

1.2.1 a. Getting the training and test sample sets from the MNIST database

(same code as PW02)

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import os
import math
from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import confusion_matrix, f1_score, classification_report
```

- We will be using the MNIST dataset from zalando <https://github.com/zalando-research/fashion-mnist>

Label	Description
0	T-shirt/top
1	Trouser
2	Pullover
3	Dress
4	Coat
5	Sandal
6	Shirt
7	Sneaker
8	Bag
9	Ankle boot

```
In [2]: # Run some setup code for this notebook.
```

```
%matplotlib inline
```

```
plt.rcParams['figure.figsize'] = (10.0, 10.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'
```

```
# This is a method to read the MNIST dataset from a ROOT directory
```

```
def load_MNIST(ROOT):
```

```
    '''load all of mnist
    training set first'''
```

```
    Xtr = []
```

```
    train = pd.read_csv(os.path.join(ROOT, 'fashion-mnist_train.csv'))
```

```
    X = np.array(train.drop('label', axis=1))
```

```
    Ytr = np.array(train['label'])
```

```
    # With this for-loop we give the data a shape of the actual image (28x28)
```

```
    # instead of the shape in file (1x784)
```

```
    for row in X:
```

```
        Xtr.append(row.reshape(28,28))
```

```
    # load test set second
```

```
    Xte = []
```

```
    test = pd.read_csv(os.path.join(ROOT, 'fashion-mnist_test.csv'))
```

```
    X = np.array(test.drop('label', axis=1))
```

```
    Yte = np.array(test['label'])
```

```
    # same reshaping
```

```
    for row in X:
```

```
        Xte.append(row.reshape(28,28))
```

```
    return np.array(Xtr), np.array(Ytr), np.array(Xte), np.array(Yte)
```

```
# Load the raw MNIST data.
```

```
mnist_dir = './data' # TODO: Fabio to update
```

```
X_train, y_train, X_test, y_test = load_MNIST(mnist_dir)
```

```
# As a sanity check, we print out the size of the training and test data.
```

```
print('Training data shape: ', X_train.shape)
```

```
print('Training labels shape: ', y_train.shape)
```

```
print('Test data shape: ', X_test.shape)
```

```
print('Test labels shape: ', y_test.shape)
```

```
# Now let's visualise some of the images
```

```
classes = ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9']
```

```
num_classes = len(classes)
```

```
samples_per_class = 7
```

```
for y, cls in enumerate(classes): # y and cls takes values from 0-9
```

```
    idxs = np.flatnonzero(y_train == y) # gets the indices of samples that corresponds
```

```

idxs = np.random.choice(idxs, samples_per_class, replace=False) # picks randomly s
for i, idx in enumerate(idxs):
    plt_idx = i * num_classes + y + 1 # determines the sub-plot index
    plt.subplot(samples_per_class, num_classes, plt_idx)
    plt.imshow(X_train[idx].astype('uint8'))
    plt.axis('off')
    if i == 0:
        plt.title(cls)
plt.show()

# Subsample the data for more efficient code execution in this exercise. We do this to
# When you will have completed the whole notebook, you can run it again on a larger (o
# and observe the difference in terms of accuracy (and speedup).
num_training = 1000
mask = range(num_training)
X_train = X_train[mask]
y_train = y_train[mask]

num_test = 500
mask = range(num_test)
X_test = X_test[mask]
y_test = y_test[mask]

# TODO: sanity check: write code to print out the size of the subsampled training and
print('Subsampled training data shape: ', X_train.shape)
print('Subsampled training labels shape: ', y_train.shape)
print('Subsampled test data shape: ', X_test.shape)
print('Subsampled test labels shape: ', y_test.shape)

```

```

Training data shape: (60000, 28, 28)
Training labels shape: (60000,)
Test data shape: (10000, 28, 28)
Test labels shape: (10000,)

```



Subsampled training data shape: (1000, 28, 28)

Subsampled training labels shape: (1000,)

Subsampled test data shape: (500, 28, 28)

Subsampled test labels shape: (500,)

1.2.2 b. Classification of digits based on raw pixel values using SVM and different kernels

```
In [3]: X_train = np.reshape(X_train, (X_train.shape[0], -1)) # when reshaping, -1 means "infer"
        X_test = np.reshape(X_test, (X_test.shape[0], -1))
```

```
In [4]: def optimize_SVM(params, datas, y, cv_folds=10, verbose=5):
        grid = GridSearchCV(SVC(cache_size=7000, max_iter=400), cv=cv_folds, n_jobs=-1,
                             param_grid=params,
                             verbose=verbose)
```

```
grid.fit(datas, y)
```

```
return grid
```

```
In [5]: rbf = optimize_SVM({
        'C': np.logspace(-2, 6, 9),
        'gamma': np.logspace(-6, 2, 9),
        'kernel': ['rbf']
    },
    X_train, y_train,
    cv_folds=7)
```

Fitting 7 folds for each of 81 candidates, totalling 567 fits

```
[Parallel(n_jobs=-1)]: Done 10 tasks      | elapsed: 8.9s
[Parallel(n_jobs=-1)]: Done 64 tasks      | elapsed: 48.7s
[Parallel(n_jobs=-1)]: Done 154 tasks     | elapsed: 2.3min
[Parallel(n_jobs=-1)]: Done 280 tasks     | elapsed: 5.2min
[Parallel(n_jobs=-1)]: Done 442 tasks     | elapsed: 9.3min
[Parallel(n_jobs=-1)]: Done 567 out of 567 | elapsed: 12.4min finished
```

```
In [7]: linear = optimize_SVM({
        'C': np.logspace(-2, 8, 11),
        'kernel': ['linear']
    },
    X_train, y_train)
```

Fitting 10 folds for each of 11 candidates, totalling 110 fits

```
[Parallel(n_jobs=-1)]: Done 10 tasks      | elapsed: 7.9s
[Parallel(n_jobs=-1)]: Done 64 tasks      | elapsed: 40.8s
[Parallel(n_jobs=-1)]: Done 110 out of 110 | elapsed: 1.2min finished
D:\ProgramData\Anaconda3\lib\site-packages\sklearn\svm\base.py:218: ConvergenceWarning: Solver
% self.max_iter, ConvergenceWarning)
```

```
In [8]: poly = optimize_SVM({
        'C': np.logspace(-1, 4, 6),
        'degree': range(1,4,1),
        'kernel': ['poly']
    },
    X_train, y_train,
    cv_folds=7)
```

Fitting 7 folds for each of 18 candidates, totalling 126 fits

```
[Parallel(n_jobs=-1)]: Done 10 tasks      | elapsed:    7.5s
[Parallel(n_jobs=-1)]: Done 64 tasks      | elapsed:   36.5s
[Parallel(n_jobs=-1)]: Done 126 out of 126 | elapsed:  1.2min finished
D:\ProgramData\Anaconda3\lib\site-packages\sklearn\svm\base.py:218: ConvergenceWarning: Solver
% self.max_iter, ConvergenceWarning)
```

```
In [9]: def print_results(system):
        predicted = system.predict(X_test)
        print(system.best_estimator_)
        print('Score: {}'.format(f1_score(y_test, predicted, average='weighted'))
        print(confusion_matrix(y_test, predicted))
        print(classification_report(predicted, y_test))

        print_results(rbf)
        print_results(linear)
        print_results(poly)
```

```
SVC(C=10.0, cache_size=7000, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma=1e-06, kernel='rbf',
    max_iter=400, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

Score: 0.7782510888501757

```
[[31  0  2  3  0  0 10  0  7  0]
 [ 0 37  0  1  0  0  1  0  1  0]
 [ 1  0 34  0 13  0  8  0  7  0]
 [ 1  0  0 43  0  0  0  0  2  0]
 [ 0  0  4  4 39  0  6  0  2  0]
 [ 0  0  0  0  0 42  0  1  8  1]
 [ 6  0  1  2  5  0 31  0  4  0]
 [ 0  0  0  0  0  1  0 45  0  6]
 [ 0  0  0  0  0  0  1  0 42  0]
 [ 0  0  0  0  0  0  0  1  1 45]]
```

	precision	recall	f1-score	support
0	0.58	0.79	0.67	39
1	0.93	1.00	0.96	37
2	0.54	0.83	0.65	41
3	0.93	0.81	0.87	53
4	0.71	0.68	0.70	57
5	0.81	0.98	0.88	43
6	0.63	0.54	0.58	57
7	0.87	0.96	0.91	47
8	0.98	0.57	0.72	74
9	0.96	0.87	0.91	52
avg / total	0.81	0.78	0.78	500

```
SVC(C=0.01, cache_size=7000, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='auto', kernel='linear',
    max_iter=400, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

Score: 0.7838199863027533

```
[[36  0  4  5  1  0  7  0  0  0]
 [ 0 40  0  0  0  0  0  0  0  0]
 [ 2  0 39  1 14  0  7  0  0  0]
 [ 4  1  0 37  3  0  1  0  0  0]
 [ 1  0  7  3 41  0  3  0  0  0]
 [ 0  0  0  0  0 45  0  4  2  1]
 [ 9  0  9  2  3  0 25  0  1  0]
 [ 0  0  0  0  0  0  0 48  0  4]
 [ 0  0  3  1  0  1  3  0 35  0]
 [ 0  0  0  0  0  0  0  1  0 46]]
```

	precision	recall	f1-score	support
0	0.68	0.69	0.69	52
1	1.00	0.98	0.99	41
2	0.62	0.63	0.62	62
3	0.80	0.76	0.78	49
4	0.75	0.66	0.70	62
5	0.87	0.98	0.92	46
6	0.51	0.54	0.53	46
7	0.92	0.91	0.91	53
8	0.81	0.92	0.86	38
9	0.98	0.90	0.94	51
avg / total	0.79	0.78	0.78	500

```
SVC(C=0.1, cache_size=7000, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=1, gamma='auto', kernel='poly',
    max_iter=400, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

Score: 0.7856829842230861

```
[[37  0  4  5  1  0  6  0  0  0]
 [ 0 40  0  0  0  0  0  0  0  0]
 [ 2  0 39  1 14  0  7  0  0  0]
 [ 4  1  0 37  3  0  1  0  0  0]
 [ 1  0  7  3 41  0  3  0  0  0]
 [ 0  0  0  0  0 45  0  4  2  1]
 [ 9  0  9  2  3  0 25  0  1  0]
 [ 0  0  0  0  0  0  0 48  0  4]
 [ 0  0  3  1  0  1  3  0 35  0]
 [ 0  0  0  0  0  0  0  1  0 46]]
```

	precision	recall	f1-score	support
0	0.70	0.70	0.70	53

1	1.00	0.98	0.99	41
2	0.62	0.63	0.62	62
3	0.80	0.76	0.78	49
4	0.75	0.66	0.70	62
5	0.87	0.98	0.92	46
6	0.51	0.56	0.53	45
7	0.92	0.91	0.91	53
8	0.81	0.92	0.86	38
9	0.98	0.90	0.94	51
avg / total	0.79	0.79	0.79	500

1.2.3 d. Analysis of the results with the best SVM classifier

- a) Which kernel and parameters were used ?

The best classifier is the polynomial kernel with the following parameters:

SVC(C=0.1, cache_size=7000, class_weight=None, coef0=0.0, decision_function_shape='ovr', degree=1, gamma='auto', kernel='poly', max_iter=400, probability=False, random_state=None, shrinking=True, tol=0.001, verbose=False)

- b) Which digit classes are the best/worse recognized against which ? Why ?

The labels 0 and 6 (T-shirts and shirts) confuse the classifier the most since they look a lot alike.

The trousers (label 1) have the best precision among all the labels. No other clothe in the set looks like trousers.

- c) What is the impact of the sizes of the training and test sets on the classification performance?

With a test set size of 500, here are the F1-score of our best classifier for different training set size:

Size of training set	Score
1000	78.5%
500	75.2%
200	73.7%